


# INTRODUCTION TO MACHINE LEARNING

## NEURAL NETWORKS I


Mingon Kang, Ph.D.  
Department of Computer Science @ UNLV

# Begin with a single Perceptron




$$a_i = g\left(\sum_j W_{j,i} a_j\right)$$


# Threshold Functions



(a) Step function




(b) Sign function




(c) Sigmoid function

# Boolean Functions and Perceptron




**AND**




**OR**


?



**(a)  $I_1$  and  $I_2$**



**(b)  $I_1$  or  $I_2$**



**(c)  $I_1$  xor  $I_2$**


# Discussion in Perceptrons

- A single perceptron works well to classify a linearly separable set of inputs
- Multi-layer perceptrons
  - ▣ found as a “solution” to represent nonlinearly separable functions (1950s)
- Many local minima – non-convex
- Research in neural networks stopped until the 70s

# XOR?

□  $\mathbf{X} = \begin{bmatrix} 0 & 0 \\ 0 & 1 \\ 1 & 0 \\ 1 & 1 \end{bmatrix}, g(x) = \max(x, 0)$

□  $\mathbf{W}_1 = [1, 1; 1, 1], \mathbf{W}_2 = [1, -2], c = [0, -1, 0]$



# Neural Network

- Feedforward Networks or Multilayer Perceptrons (MLPs)

- E.g., for a classifier

- $y = f^*(x; \theta)$  maps an input  $x$  to a category  $y$

- A feedforward network defines a mapping  $y = f(x; \theta)$  and learns the optimal parameters  $\theta$ .

# Neural Network

## □ Called Networks

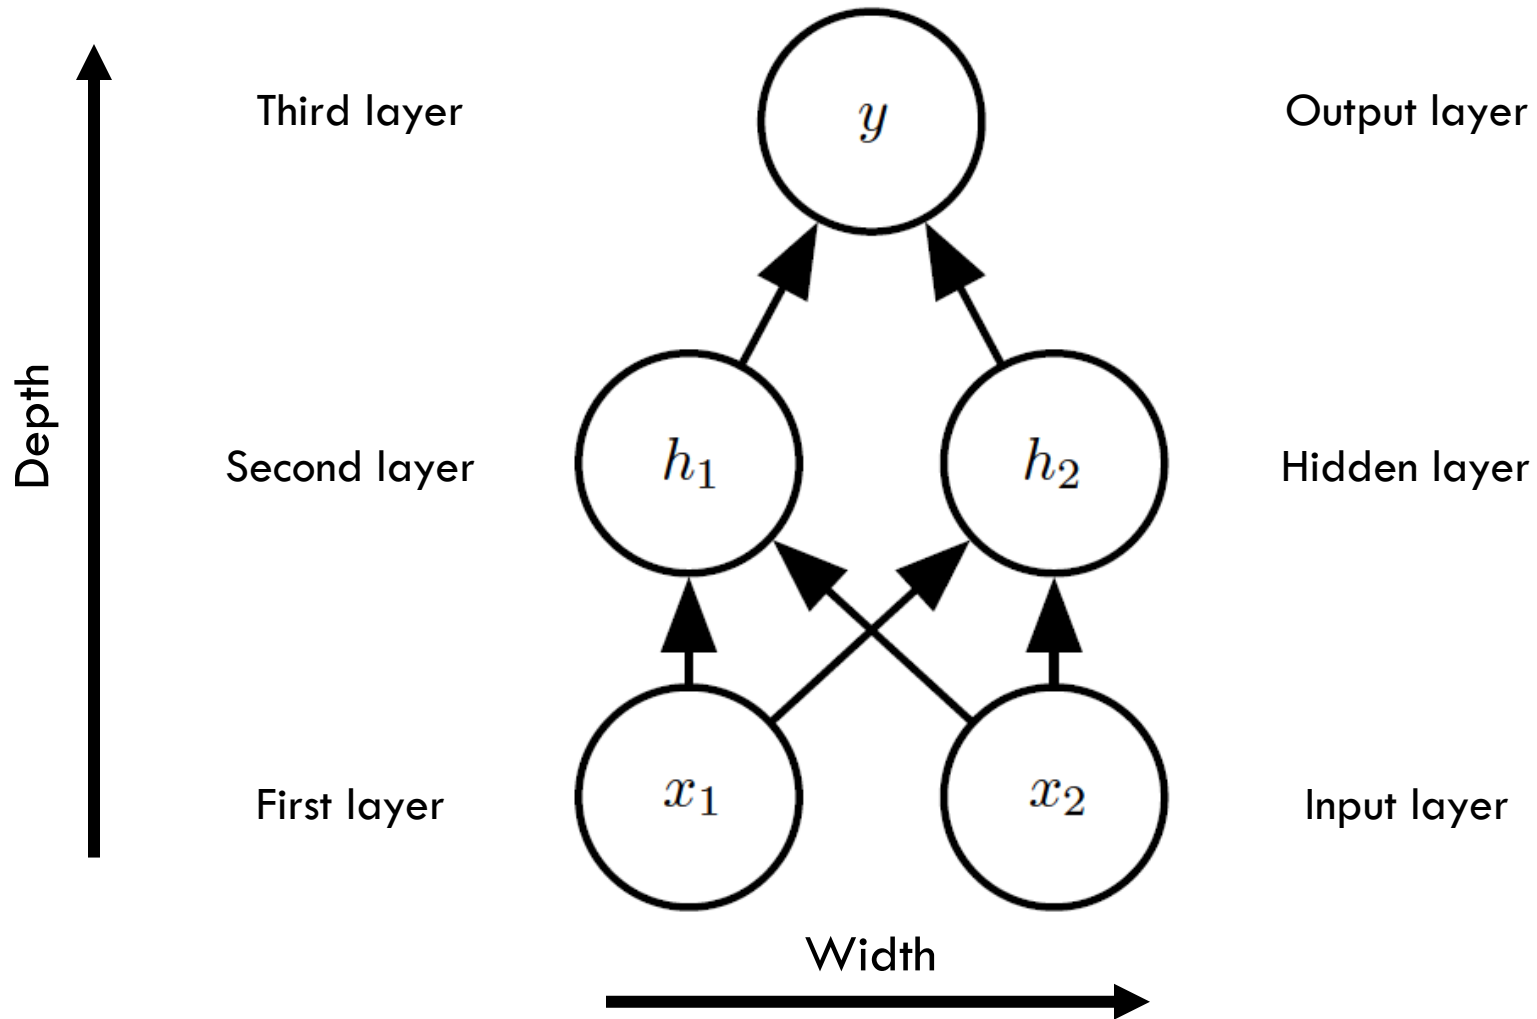
- ▣ Involve many various functions
- ▣ Associated with a directed acyclic graph that represent how the functions are composed together
- ▣ Mostly, chain structures

- $$f(x) = f^{(3)} \left( f^{(2)} \left( f^{(1)}(x) \right) \right)$$

$f^{(i)}$ : the i-th layer in the network




# Neural Network



# Design of a NN model

- Architecture of the network
  - ▣ How many layers
  - ▣ How these layers are connected to each other
  - ▣ How many units are in each layer
- Activation function: to compute the hidden layer values
- Cost function: to optimize the model
- Optimizer: how to optimize the model

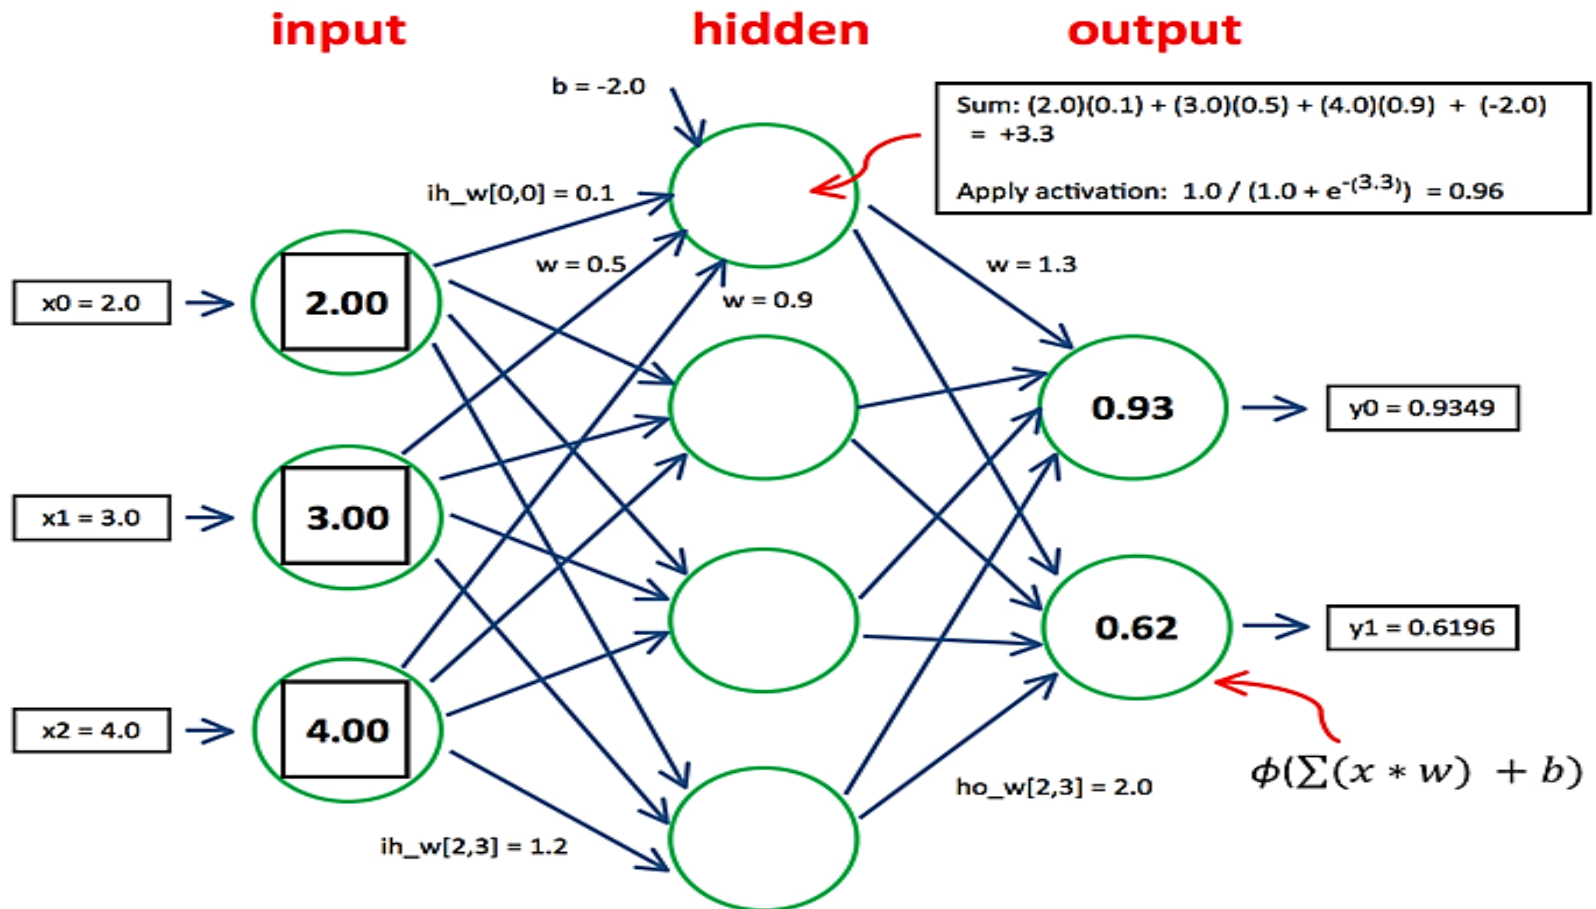
# Weights and Bias



# Weighted Input and Activation of a Neuron

- Weighted Input of the  $j^{th}$  neuron of the  $l^{th}$  layer is  $z_j^l$ :
  - $z_j^l = \sum_k \left( (w_{jk}^l a_k^{l-1}) + b_j^l \right)$
- Activation from the  $j^{th}$  neuron of the  $l^{th}$  layer is  $a_j^l$ :
  - $a_j^l = f(z_j^l)$  , where  $f$  is the activation function
  - $a_j^l = f \left( \sum_k \left( (w_{jk}^l a_k^{l-1}) + b_j^l \right) \right)$
- In matrix notation, the activation becomes:
  - $a^l = f(w^l a^{l-1} + b^l)$

# Example



# Activation Functions


- Sigmoid Activation:  $a_j^l = \sigma(z_j^l) = \frac{1}{1 + e^{-z_j^l}}$
- Softmax Activation:  $a_j^l = \frac{e^{z_j^l}}{\sum_k e^{z_k^l}}$
- Tanh Activation:  $\tanh(z_j^l) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$
- Rectified Linear Activation:  $\max(0, z_j^l)$ , maximum of 0 or  $z_j^l$

# Activation Functions

## Activation Functions


### Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$




### tanh

$$\tanh(x)$$




### ReLU

$$\max(0, x)$$



### Leaky ReLU

$$\max(0.1x, x)$$




### Maxout










$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

### ELU

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$




# Activation Functions

Name	Plot	Equation	Derivative
Identity		$f(x) = x$	$f'(x) = 1$
Binary step		$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} 0 & \text{for } x \neq 0 \\ ? & \text{for } x = 0 \end{cases}$
Logistic (a.k.a Soft step)		$f(x) = \frac{1}{1 + e^{-x}}$	$f'(x) = f(x)(1 - f(x))$
TanH		$f(x) = \tanh(x) = \frac{2}{1 + e^{-2x}} - 1$	$f'(x) = 1 - f(x)^2$
ArcTan		$f(x) = \tan^{-1}(x)$	$f'(x) = \frac{1}{x^2 + 1}$
Rectified Linear Unit (ReLU)		$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} 0 & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$
Parametric Rectified Linear Unit (PReLU) [2]		$f(x) = \begin{cases} \alpha x & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} \alpha & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$
Exponential Linear Unit (ELU) [3]		$f(x) = \begin{cases} \alpha(e^x - 1) & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} f(x) + \alpha & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$
SoftPlus		$f(x) = \log_e(1 + e^x)$	$f'(x) = \frac{1}{1 + e^{-x}}$



# Universal Approximator Theorem

- One hidden layer may be enough to represent (not learn) an approximation of any function to an arbitrary degree of accuracy
- So why deeper?
  - ▣ Shallow net may need (exponentially) more width
  - ▣ Shallow net may overfit more



# Cost Functions

- Quadratic Cost:  $C = \frac{1}{2n} \sum_x ||y(x) - a^L(x)||^2$
- Binary Cross-Entropy Cost:  $C = -\frac{1}{n} \sum_x (y(x) \times \log_e(a^L(x)) + (1 - y) \times \log_e(1 - a^L(x)))$
- Negative Log-likelihood Cost:  $C = -\log_e(a_y^L)$
- A Cost function must satisfy the following two conditions (for backpropagation):
  - The Cost function  $C$  should be calculated as an average over the cost functions  $C_x$  for individual training examples.
  - The cost functions for the individual training examples  $C_x$  and consequently the Cost  $C$  function must be a function of the outputs of the neural network.

# Examples of Neg Log-likelihood Cost

- Given the posterior probability and the ground truth:
  - ▣ A set of output probabilities: e.g. [0.1, 0.3, 0.5, 0.1]
  - ▣ Ground truth: e.g., [0, 0, 0, 1]
- Likelihood
  - ▣  $0*0.1 + 0*0.3 + 0*0.5 + 1*0.1 = 0.1$
  - ▣ NLL:  $-\ln(0.1) = 2.3$
- If ground truth is [0, 0, 1, 0]
  - ▣  $0*0.1 + 0*0.3 + 1*0.5 + 0*0.1 = 0.5$
  - ▣ NLL:  $-\ln(0.5) = 0.69$

# Output Types

Output Type	Output Distribution	Output Layer	Cost Function
Binary	Bernoulli	Sigmoid	Binary cross-entropy
Discrete	Multinoulli	Softmax	Discrete cross-entropy
Continuous	Gaussian	Linear	Gaussian cross-entropy (MSE)
Continuous	Mixture of Gaussian	Mixture Density	Cross-entropy
Continuous	Arbitrary	See part III: GAN, VAE, FVBN	Various

# Discussion

- Construct a neural network for MNIST data
  - ▣ 28\*28 pixel images and 10 labels
- Construct a neural network for a binary classification problem
  - ▣ 10 features and 2 labels
- Construct a neural network for predicting tomorrow's temperature in degree
  - ▣ 10 variables