

# Spark

Fast, Interactive, Language-Integrated  
Cluster Computing

Matei Zaharia, Mosharaf Chowdhury, Tathagata Das,  
Ankur Dave, Justin Ma, Murphy McCauley, Michael Franklin,  
Scott Shenker, Ion Stoica

[www.spark-project.org](http://www.spark-project.org)



# Project Goals

Extend the MapReduce model to better support two common classes of analytics apps:


- » **Iterative** algorithms (machine learning, graphs)
- » **Interactive** data mining

Enhance programmability:

- » Integrate into Scala programming language
- » Allow interactive use from Scala interpreter

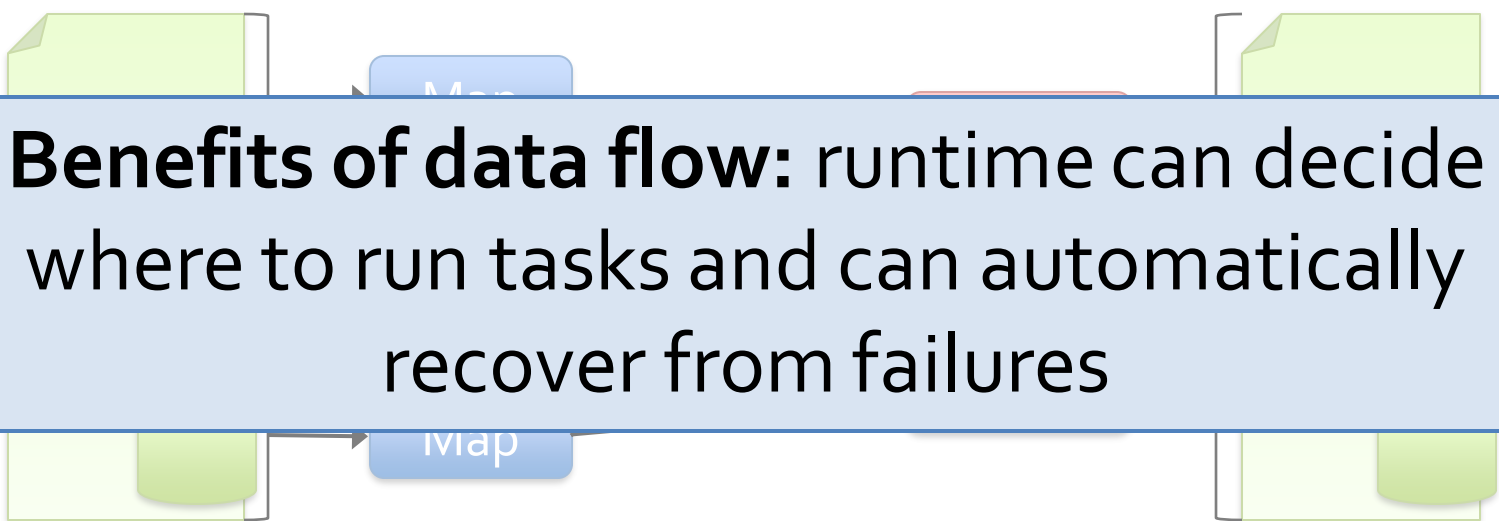
# Motivation

Most current cluster programming models are based on *acyclic data flow* from stable storage to stable storage



# Motivation

Most current cluster programming models are based on *acyclic data flow* from stable storage to stable storage



The diagram shows a data flow graph. At the top, a light green document icon is connected to a blue rounded rectangle labeled 'Map'. Below this, a light green document icon is connected to a blue rounded rectangle labeled 'MapReduce'. At the bottom, a light green document icon is connected to a blue rounded rectangle labeled 'Map'. Arrows indicate a flow from the top 'Map' node to the bottom 'Map' node, and from the bottom 'Map' node to the 'MapReduce' node. A large blue rounded rectangle with a thin border is overlaid on the diagram, containing the text 'Benefits of data flow: runtime can decide where to run tasks and can automatically recover from failures'.

**Benefits of data flow:** runtime can decide where to run tasks and can automatically recover from failures

# Motivation

Acyclic data flow is inefficient for applications that repeatedly reuse a *working set* of data:

- » **Iterative** algorithms (machine learning, graphs)
- » **Interactive** data mining tools (R, Excel, Python)

With current frameworks, apps reload data from stable storage on each query

# **Solution: Resilient Distributed Datasets (RDDs)**

Allow apps to keep working sets in memory for efficient reuse

Retain the attractive properties of MapReduce  
» Fault tolerance, data locality, scalability

Support a wide range of applications

# RDD

- Resilient Distributed Datasets: A Fault-Tolerant Abstraction for In-Memory Cluster Computing (2012)
- Most machine learning algorithms require iterative computation.
- The iterations on MapReduce cause big overhead between Map and Reduce
  - Data replication
  - Disk I/O
  - Serialization

# RDD

- The iterations are computationally expensive since Hadoop uses HDFS for sharing data
- HDFS causes frequent file I/O → Slow
- Solutions
  - Reduce uses of file I/O
  - Use RAM



# RDD


- Using RAM is much more efficient
  - However, how to handle fault-tolerant?
  - Need to load the data again into memory?
- Instead update data in RAM, make all data in RAM as read-only.

# RDD

- Designed by Lineage and Directed Acyclic Graph (DAG)
  - RDD records all history of the process of the data
  - Fault-tolerant happens, RDD checks the lineage of the data and roll back → Fast recovery
  - All data is stored as DAG, so efficient.

# RDD

- Lineage and DAG



# Outline

Spark programming model

Implementation

Demo

User applications

# Programming Model

## Resilient distributed datasets (RDDs)

- » Immutable, partitioned collections of objects
- » Created through parallel *transformations* (map, filter, groupBy, join, ...) on data in stable storage
- » Can be *cached* for efficient reuse

## *Actions* on RDDs

- » Count, reduce, collect, save, ...

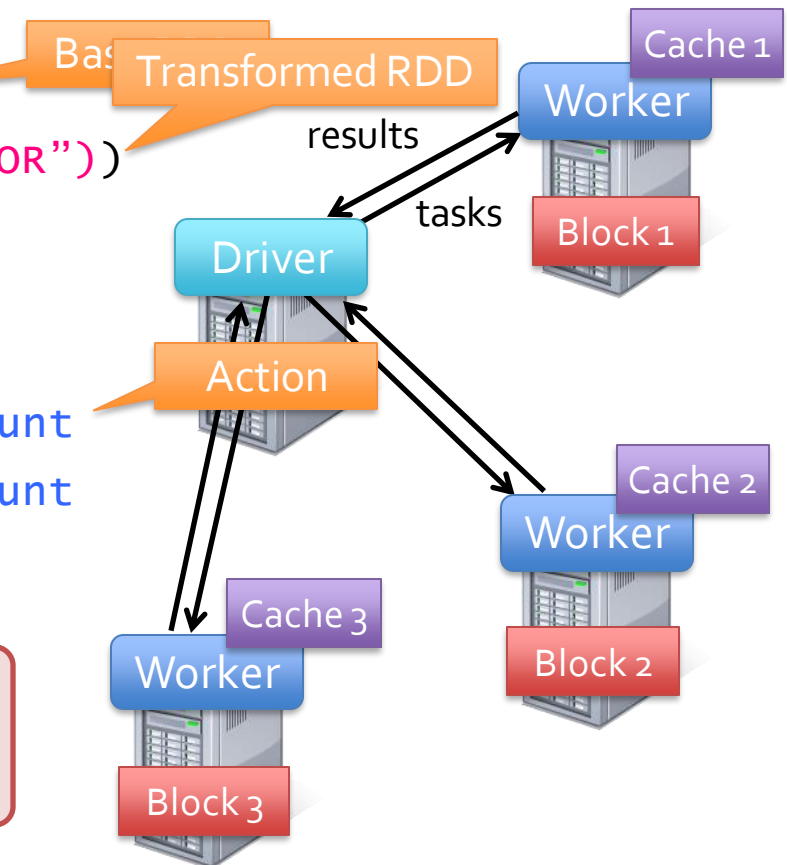
# Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns

```
lines = spark.textFile("hdfs://...")
errors = lines.filter(_.startsWith("ERROR"))
messages = errors.map(_.split('\t')(2))
cachedMsgs = messages.cache()

cachedMsgs.filter(_.contains("foo")).count
cachedMsgs.filter(_.contains("bar")).count
. . .
```

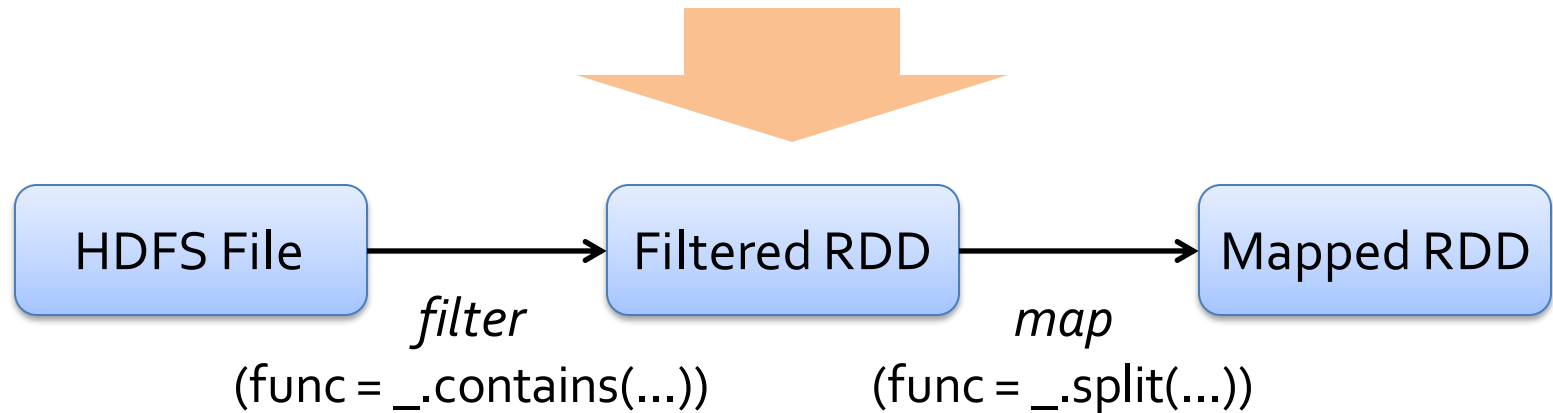
**Result:** scaled to 1 TB data in 5-7 sec  
(vs 170 sec for on-disk data)



# RDD Fault Tolerance

RDDs maintain *lineage* information that can be used to reconstruct lost partitions

Ex: `messages = textFile(...).filter(_.startsWith("ERROR")).map(_.split('\t')(2))`



# Word Count

Use a few transformations to build a dataset of (String, int) pairs called counts and then save it to a file.

```
val textFile = sc.textFile("hdfs://...")
val counts = textFile.flatMap(line => line.split(" "))
                      .map(word => (word, 1))
                      .reduceByKey(_ + _)
counts.saveAsTextFile("hdfs://...")
```

The map operation produces one output value for each input value, whereas the flatMap operation produces an arbitrary number (zero or more) values for each input value




# Pi estimation

This code estimates  $\pi$  by "throwing darts" at a circle. We pick random points in the unit square ((0, 0) to (1,1)) and see how many fall in the unit circle. The fraction should be  $\pi / 4$ , so we use this to get our estimate.

```
val count = sc.parallelize(1 to NUM_SAMPLES).map{i =>
  val x = Math.random()
  val y = Math.random()
  if (x*x + y*y < 1) 1 else 0
}.reduce(_ + _)
println("Pi is roughly " + 4.0 * count / NUM_SAMPLES)
```

# Logistic Regression Performance



# Spark Applications

In-memory data mining on Hive data (Conviva)

Predictive analytics (Quantifind)

City traffic prediction (Mobile Millennium)

Twitter spam classification (Monarch)

Collaborative filtering via matrix factorization

...

# Interactive Spark

Modified Scala interpreter to allow Spark to be used interactively from the command line

Required two changes:

- » Modified wrapper code generation so that each line typed has references to objects for its dependencies
- » Distribute generated classes over the network

# Conclusion

Spark provides a simple, efficient, and powerful programming model for a wide range of apps

Download our open source release:

[www.spark-project.org](http://www.spark-project.org)

matei@berkeley.edu

# Related Work

## DryadLINQ, FlumeJava

- » Similar “distributed collection” API, but cannot reuse datasets efficiently *across* queries

## Relational databases

- » Lineage/provenance, logical logging, materialized views

## GraphLab, Piccolo, BigTable, RAMCloud

- » Fine-grained writes similar to distributed shared memory


## Iterative MapReduce (e.g. Twister, HaLoop)

- » Implicit data sharing for a fixed computation pattern

## Caching systems (e.g. Nectar)

- » Store data in files, no explicit control over what is cached

# Behavior with Not Enough RAM



# Fault Recovery Results

